

10 leading AI model providers and models

Below is a consolidated list of **10 leading AI model providers and models released within the last 12 months (July 2025 – July 2026)**. Pricing and context windows change frequently, so I've used the latest publicly available information where available.

Model Provider	Model	Release Date	Knowledge Cutoff	Token Cost (Input / Output)	Context Window	Type	When to Use
OpenAI	GPT-5.6	Jul 2026	Not publicly disclosed	Not yet publicly disclosed	Not yet disclosed	Multimodal Reasoning	General-purpose AI, coding, agents, enterprise workflows (Axios)
Anthropic	Claude Opus 4	May 22, 2025	Dynamic / training cutoff not disclosed	\$15 / \$75 per 1M tokens	200K (up to 1M in some tiers)	Advanced Reasoning & Coding	Complex coding, long-running agents, research workflows (Anthropic)
Anthropic	Claude Sonnet 4	May 22, 2025	Dynamic / training cutoff not disclosed	\$3 / \$15 per 1M tokens	200K (up to 1M in some tiers)	Balanced Reasoning	Everyday coding, business applications, cost-efficient AI assistants (Anthropic)
Google Gemini	Gemini 2.5 Pro	Mar 25, 2025	Not publicly disclosed	Varies by platform	Up to 1M tokens	Multimodal Reasoning	Large document analysis, coding, research, multimodal tasks (Gemini)
xAI	Grok 4	Jul 9, 2025	Real-time web access	~\$3 / \$15 per 1M tokens (API launch pricing)	256K+	Reasoning Model	Deep reasoning, math, science, real-time information retrieval (SpaceXAI)
xAI	Grok 4.3	Jun 2026	Real-time web access	\$1.25 / \$2.50 per 1M tokens	1M tokens	Agentic Multimodal	High-volume production workloads, tool calling, enterprise AI (SpaceXAI Docs)

Model Provider	Model	Release Date	Knowledge Cutoff	Token Cost (Input / Output)	Context Window	Type	When to Use
xAI	Grok 4.20 Reasoning	2026	Real-time web access	\$1.25 / \$2.50 per 1M tokens	1M tokens	Reasoning Specialist	Long-context reasoning, tool-using agents, knowledge-intensive tasks (SpaceXAI Docs)
Meta AI	Llama 4 Scout	Apr 2025	Dec 2024 (approx.)	Open-weight (self-hosted)	10M tokens	Open-weight Multimodal	Massive context workloads, enterprise self-hosting
Meta AI	Llama 4 Maverick	Apr 2025	Dec 2024 (approx.)	Open-weight (self-hosted)	1M tokens	Open-weight General Purpose	Production AI assistants, fine-tuning, enterprise deployment
Mistral AI	Magistral Medium	Jun 2026	Not disclosed	Usage-based API pricing	128K+	Reasoning Model	Enterprise reasoning, multilingual business applications

Recommended Models by Use Case

Use Case	Best Model
Software Development	Claude Opus 4
Cost-Efficient Coding	Claude Sonnet 4
Massive Document Analysis	Gemini 2.5 Pro
Longest Context Window	Llama 4 Scout (10M)
Real-Time Information	Grok 4 / Grok 4.3
Enterprise AI Agents	GPT-5.6
Open-Source Deployment	Llama 4 Maverick
Research & Complex Reasoning	Claude Opus 4, Grok 4
Multimodal Tasks	GPT-5.6, Gemini 2.5 Pro
Budget Production APIs	Grok 4.3

Key Trends in the Last 12 Months

1. **Context windows expanded dramatically** from 128K–200K to **1M+ tokens**.
2. **Reasoning models** became the default frontier architecture.
3. **Agentic tool use** (web search, coding, file systems) is now standard.
4. **Multimodal capabilities** (text, image, audio, video) are increasingly integrated.
5. **Cost efficiency improved**, with some frontier models now below \$3 per million input tokens.
([Anthropic](#))